

CENG499

INTRODUCTION TO MACHINE LEARNING

THE3

REPORT

Erdem
Şamlıoğlu
2448843

1 Part1

Gain Ratio Criterion

Selected features in order and with their Gain Ratio values.

- **milk** (Gain: 1.0)
- **fins** (Gain: 1.0000000000000002)
- **feathers** (Gain: 0.9999999999999999)
- **backbone** (Gain: 0.9999999999999999)
- **airborne** (Gain: 0.6073508754546144)
- **predator** (Gain: 0.3448660840774643)
- **legs** (Gain: 1.0)
- **tail** (Gain: 0.6073508754546144)
- **aquatic** (Gain: 0.3448660840774643)
- **eggs** (Gain: 1.0)

Information Gain Criterion

Selected features in order and with their Information Gain values.

- **legs** (Gain: 1.3630469031539394)
- **fins** (Gain: 0.8865408928220899)

- **toothed** (Gain: 0.9852281360342515)
- **eggs** (Gain: 0.6962122601251458)
- **hair** (Gain: 0.8256265261578954)
- **hair** (Gain: 0.6892019851173654)
- **aquatic** (Gain: 0.8631205685666308)
- **toothed** (Gain: 0.7219280948873623)
- **aquatic** (Gain: 0.7219280948873623)

2 Part2

2.1 Dataset 1

For Dataset 1, I have used 4 different configurations(C: 1 and 10, kernel:linear and rbf)




Figure 1: SVM with $C = 1$, Kernel=Linear

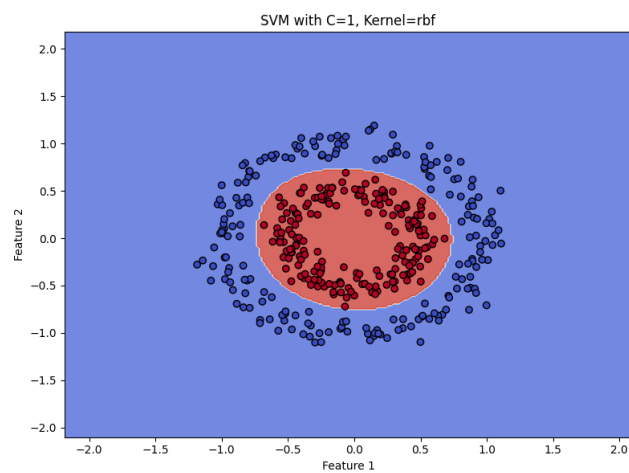


Figure 2: SVM with $C = 1$, Kernel=RBF




Figure 3: SVM with $C = 10$, Kernel=Linear




Figure 4: SVM with $C = 10$, Kernel=RBF

2.2 Dataset 2

Mean Accuracy results of the 5 repetitions of 10 fold cross validation with 4 different configurations:

| Configuration | C | Kernel | Mean Accuracy | 95% Confidence Interval |
|---------------|----|--------|---------------|-------------------------|
| 1 | 1 | Linear | 0.748 | [0.737, 0.759] |
| 2 | 1 | RBF | 0.763 | [0.754, 0.772] |
| 3 | 10 | Linear | 0.748 | [0.737, 0.759] |
| 4 | 10 | RBF | 0.753 | [0.743, 0.764] |

Table 1: Configurations' Mean Accuracy Results

The best configuration with the highest mean accuracy is Kernel:RBF and C:1. Considering the confidence intervals are overlapping and mean accuracies are really close, all these four configurations are similar in performance but as it is stated RBF kernel and C:1 is slightly beter in mean accuracy metric.

3 Part3

| Algorithm | Most-Chosen Hyperparameters | Times Chosen | Mean F1 | 95% CI |
|-------------------|---|--------------|---------|------------------|
| KNN | {n_neighbors=3, metric=manhattan} | 15/15 | 0.8552 | [0.8510, 0.8594] |
| SVM | {C=1.0, kernel=rbf} | 15/15 | 0.8279 | [0.8223, 0.8335] |
| Decision Tree | {max_depth=10, criterion=entropy} | 8/15 | 0.7982 | [0.7894, 0.8070] |
| Decision Tree | {max_depth=10, criterion=gini} | 7/15 | | |
| Random Forest | {max_depth=10, n_estimators=100} | 15/15 | 0.8633 | [0.8578, 0.8688] |
| MLP | {hidden_layer_sizes=(100,), alpha=0.001} | 6/15 | 0.8518 | [0.8480, 0.8556] |
| MLP | {hidden_layer_sizes=(100,), alpha=0.0001} | 9/15 | 0.8518 | |
| Gradient Boosting | {n_estimators=100, learning_rate=0.1} | 15/15 | 0.8482 | [0.8422, 0.8543] |

Table 2: Nested Cross Validation F1 results and which hyperparameters are selected.

Hyperparameter configurations used:

- **Hyperparameter Configurations:**

- KNN: {n_neighbors=[3,5], metric=[*euclidean*, *manhattan*]}
- SVM: {C=[0.1,1.0], kernel=[*linear*, *rbf*]}
- Decision Tree: {max_depth=[5,10], criterion=[*gini*, *entropy*]}
- Random Forest: {n_estimators=[50,100], max_depth=[5,10]}
- MLP: {hidden_layer_sizes=[(50,),(100,)], alpha=[1e-3,1e-4]}
- Gradient Boosting: {n_estimators=[50,100], learning_rate=[0.01,0.1]}

I used `random_state` to eliminate randomization.

As shown, Random Forest obtains the highest mean F1 with 0.8633 with 95% CI [0.8578, 0.8688], also KNN with 0.8552 and MLP with 0.8518 are close to that. For certain algorithms, SVM, KNN, Random Forest, chose the same parameter set for all folds, showing a strong hyperparameter combination (among the ones I have tested) among others.

Decision Tree chose different criteria, gini and entropy, time to time and ended up with a lower overall F1.

Gradient Boosting showed good performance but it took much longer in runtime compared to other methods. (In my PC it took 12-13 hours to run whole methods, Gradient Boosting took 9-10 hours itself. Later I tried it in Google Colab, It took 2-2.5 hours and Gradient Boosting took around 1.5 hours.)

In conclusion, for this ECG dataset, Random Forest yields the best average F1 among the tested methods and hyperparameter grids.